

The State of Haskell in Ethereum

author: Martin Allen, FOAM

email: martin@foam.space

date: June 27, 2018

Introduction

- What is Ethereum? Where does it fit?
- What makes Ethereum development "difficult"?
- What off the shelf haskell libraries can help?

Blockchain vs Cryptocurrency

- Cryptocurrency ~ programmable money
 - access control
 - specified utility
 - scarcity
- Blockchain ~ řŤ
 - process digital assets¹
 - cryptographic²
 - distributed

[1] This is currently the primary use case, but all work is leading towards more general use cases, e.g. governance.

[2] In the sense of cryptographic hash functions.

Blockchains in Production

- Existing: Bitcoin, Ethereum
- Nascent: Kadena, Cardano SL, RChain

Ethereum Today

- Only public, permissionless, 2nd generation blockchain in production
- Largest developer community
- Will be here for a while...

Example Applications

- Tokenized Assets
 - Digital Art -- [SuperRare](#), [CryptoKitties](#)
 - Intellectual Property / Licensing
- Digital Registries
 - Ad Tech -- [ad chain](#)
 - Geospatial Data -- [FOAM](#)

Data Access and Streaming

(An Ethereum Node's point of view)

The Problem

- Node's primary concern is block production, propagation, and verification
- Query performance and expression are not primary concerns.

Solution: Wherever applicable, use production grade databases as a history cache (e.g. Postgres, Elasticsearch).

Access Model

Libraries give the appearance of a high level interface:

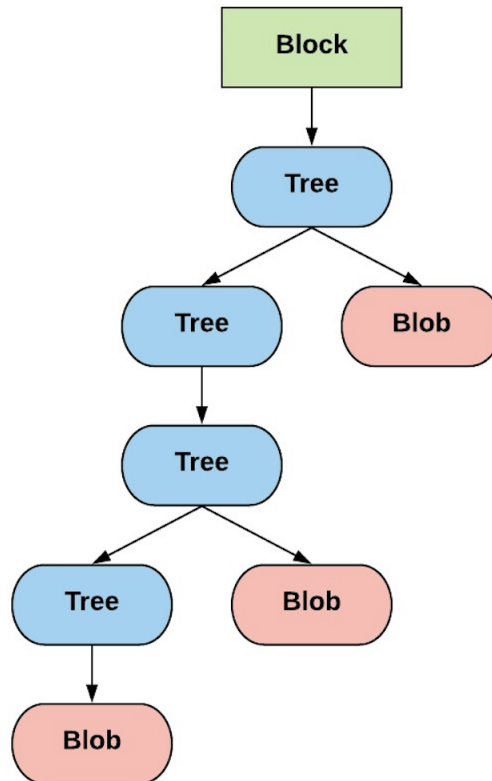
```
-- | Get the token balance for a user
balanceOf
  :: ContractAddress
  -> BlockNumber
  -> UserAddress
  -> Web3 TokenBalance
balanceOf = ...
```

Internally, the low level interface is key value storage:

```
getValue :: ByteString -> Maybe ByteString
```

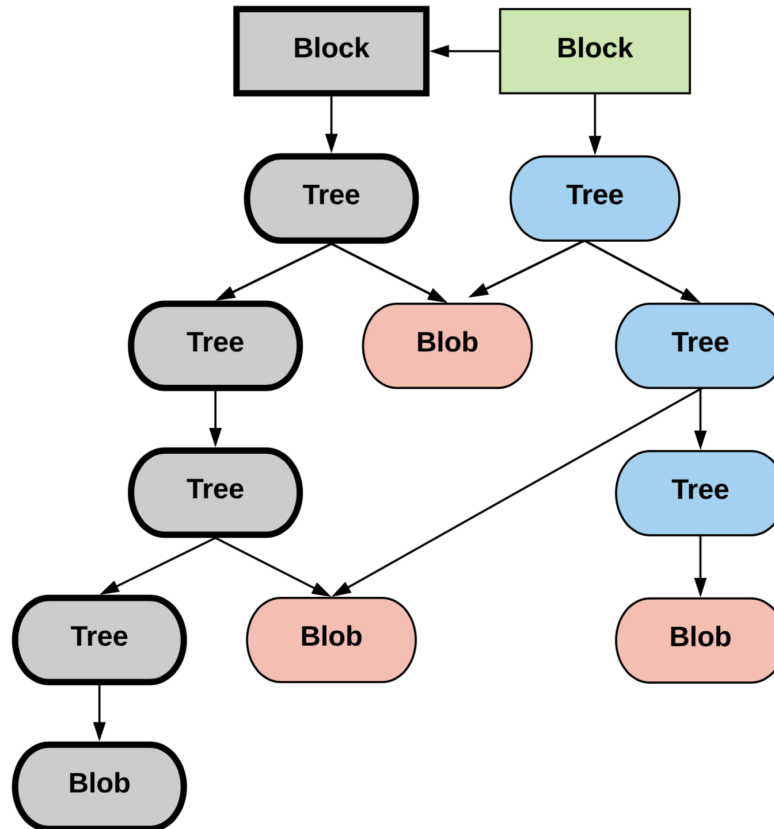
Ethereum Storage

- Internal storage model is cryptographic (Merkle) tree, not meant for Enterprise data access.
- Each block contains a different index into this tree, indicating the current state.



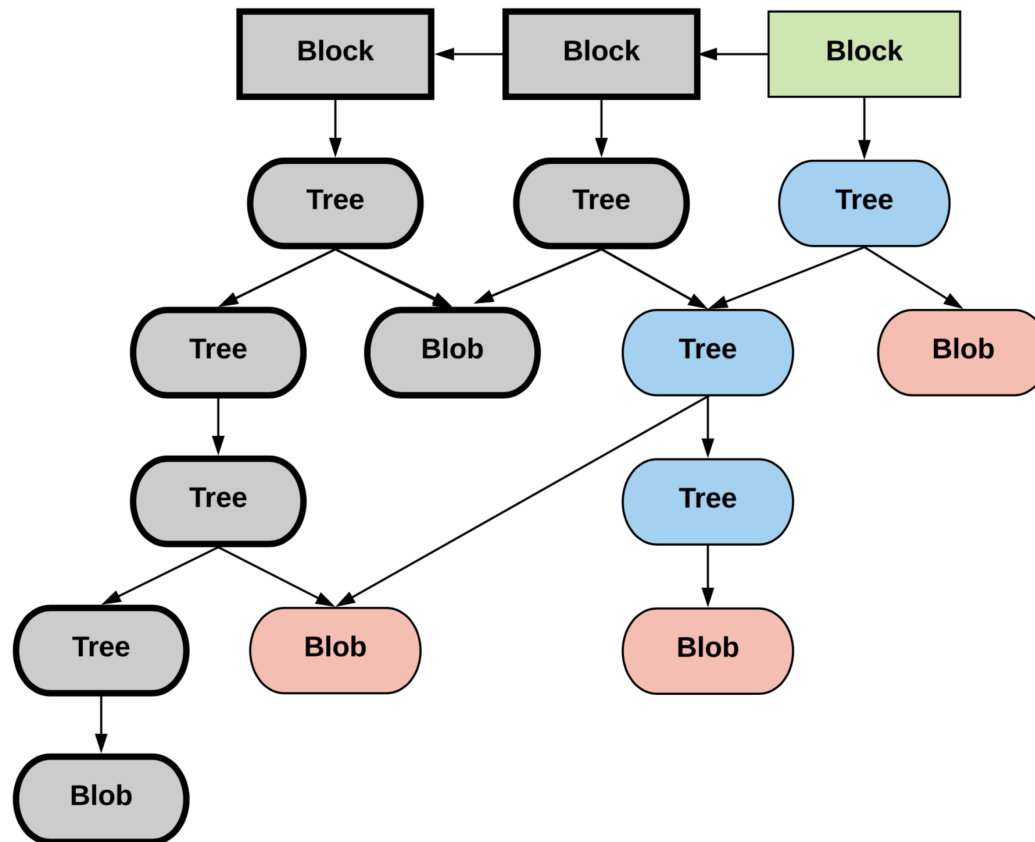
Ethereum Block Progression

- Storage is never overwritten.
- Block progression gives a series of indices into storage, i.e. Ethereum is a *time-indexed database*.



Ethereum Block Progression

- Smart Contract state changes *could* be inferred via tree diffs.
- *Events* are the mechanism to observe and react to state changes for external applications.



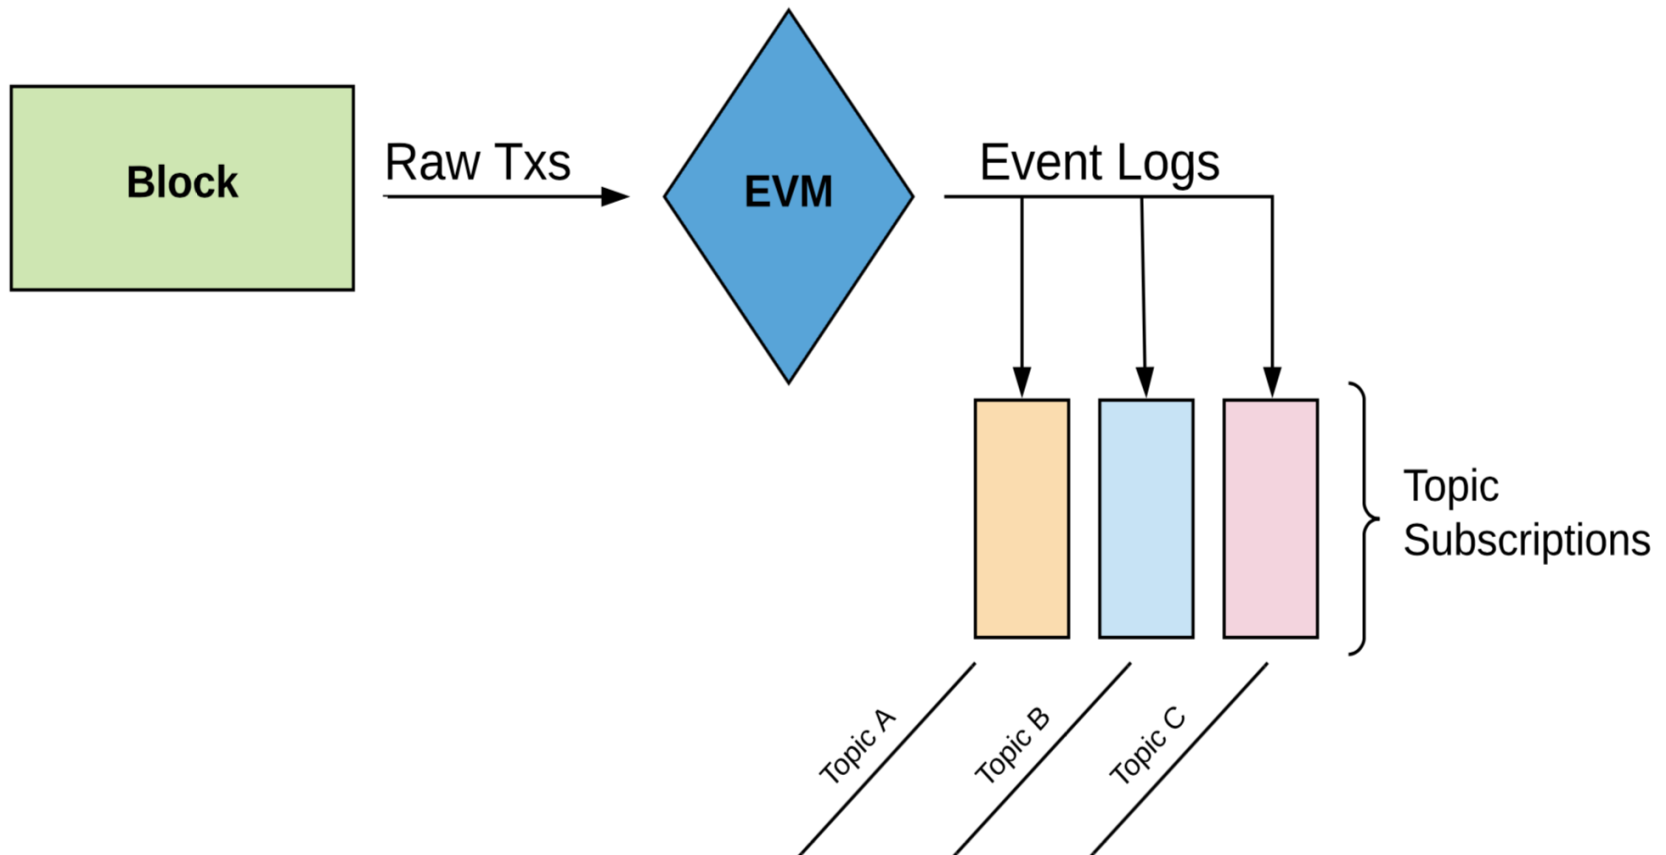
What Are These *Event Logs* Precisely? (Example)

Example:

- Alice transfers Bob a token.
- Their balances in contract storage change (new state tree index).
- This act of exchange `Transfer(Alice, Bob, 1)` is not stored in the contract-- it isn't even accessible.
- Outside world notified via contract's *Transfer* topic.

What Are These *Event Logs* Precisely? (Illustration)

```
processBlock  
:: Block  
-> StateT StateTree (Writer EventLog) ()
```



Functional Tools

Minimal web3 Library Requirements

- stream relevant Smart Contract updates, i.e. *events*
- access state history (e.g. indexing, auditing, replaying)
- submit transactions to update the blockchain

Core Functional Ethereum Libraries

- **hs-web3**
- `purescript-web3` - similar design to `hs-web3`.
- `chanterelle` - smart contract build/deploy/test.

Interacting with Smart Contracts

- Ethereum smart contracts expose an interface
- Example: ERC20 standard

The ERC20 Token Standard Interface

Following is an interface contract declaring the required functions and events to meet the ERC20 standard:

```
1 // -----
2 // ERC Token Standard #20 Interface
3 // https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20-token-standard.md
4 // -----
5 contract ERC20Interface {
6     function totalSupply() public constant returns (uint);
7     function balanceOf(address tokenOwner) public constant returns (uint balance);
8     function allowance(address tokenOwner, address spender) public constant returns (uint remaining);
9     function transfer(address to, uint tokens) public returns (bool success);
10    function approve(address spender, uint tokens) public returns (bool success);
11    function transferFrom(address from, address to, uint tokens) public returns (bool success);
12
13    event Transfer(address indexed from, address indexed to, uint tokens);
14    event Approval(address indexed tokenOwner, address indexed spender, uint tokens);
15 }
```

Interacting with Smart Contracts (ABI)

- Compilation artifact called ABI (application binary interface)
- JSON object specifying interface:
 - function calls
 - events / topics
- minimum data about contract required to generate FFI in any language (ps, haskell).

Interacting with Smart Contracts (QuasiQuoter)

```
[abiFrom|".abis/ERC20.json"]
```

```
...  
-- | Transfer a quantity of tokens to an address.  
transfer :: Call -> Address -> UIntN 256 -> Web3 Hash  
  
-- | Represents a 'Transfer' event log.  
data Transfer =  
    Transfer { transferFrom :: Address  
                , transferTo  :: Address  
                , transferValue :: UIntN 256  
                } deriving (Show, Eq, Ord, Generic)  
  
...
```

Interacting with Smart Contracts (Encodings)

- Ethereum has its own binary encoding schema (ABI-codecs)
 - transactions are serialized *closures* (function + args)
 - events / topics
- codecs derived generically (`generics-sop`)
- all other required instances declared in QQ

Conclusion:

All datatypes and FFI needed to interact with *any* Smart Contract can be generated from the QuasiQuoter and used natively in any application.

Streaming Logs

Machines - *Ceci n'est pas une pipe*

- similar libraries:
 - *conduit*
 - *purescript-coroutines*
 - *pipes*
- Idea: build up data processing machines¹, compose them, join them, split them.
- Useful for cases when streaming from one IO source to another with intermediate processing phases.

1. Usually build something called a *Plan*, which is a DSL describing what the machine should do and how to terminate.

Analogous Problem

You want to stream all photos uploaded to some api by a user.

- API supports `GET /photos/user_id` .
- supports ws notifications for new photos from a `user_id` after time `t` .

1. Break up the http requests into intervals, use a cursor.
2. Whenever you've caught up to the current time, start the ws.

This is what the SuperRare monitor is doing with Ethereum + IPFS.

Filter Type Description

We send out a description of what we want to monitor.

```
-- | Filter has a phantom type for the specific event.
data Filter e =
  Filter { address      :: [Address]
        -- ^ contract addresses for event origin.
        , fromBlock    :: ChainCursor
        , toBlock      :: ChainCursor
        , topics       :: [Maybe ByteString]
        -- ^ Bloom filter to narrow this event filter.
        }
```

Each match that comes back looks like this:

```
data FilterChange e =
  FilterChange { rawChange :: Change
               -- ^ event metadata (e.g. tx hash)
               , event      :: e
               -- ^ particular event (e.g. Transfer)
               }
```

Filter Machines (slightly simplified)

There's an internal distinction between folding over past blocks and subscribing to current events (think of example problem).

Folding over past events:

```
-- | Take a filter and split it into smaller intervals.  
filterStream  
  :: FilterStreamState e  
  -> MachineT Web3 k (Filter e)  
  
-- | Stream past events until you are caught up to  
-- the end of the filter's interval.  
playLogs  
  :: DecodeEvent i ni e  
  => FilterStreamState e  
  -> MachineT Web3 k (FilterChange e)
```

Weirdly, all web3 machines are polymorphic in `k`

Filter Machines (slightly simplified)

Subscribing to current events:

```
-- | Poll the filter until a given block number
-- TODO: support ws
pollFilter
  :: forall i n e k .
    DecodeEvent i n e
=> Filter e
-> MachineT Web3 k (FilterChange e)
```

Reduce a machine by running it:

```
data EventAction = Continue | Terminate

-- | run while 'mapM_'ing with the given handler.
-- return conditions of termination.
reduceEventStream
  :: Monad m
=> MachineT m k (FilterChange e)
-> (e -> ReaderT Change m EventAction)
-> m (Maybe (EventAction, BlockNumber))
```

High level functions

There are also high level functions that don't require you to interface with machines.

```
-- | Subscribe to the given filter, processing a number  
-- | of blocks at a time with the handler until  
-- | possibly transitioning to polling or exiting.
```

eventMany

```
:: DecodeEvent i ni e  
=> Filter e  
-> WindowSize  
-> (e -> ReaderT Change Web3 EventAction)  
-> Web3 ()
```

Ironically, I usually use this to feed a `Conduit` in all the data processing work that I do.

Querying Data

Haxl Interlude



What is Haxl?

- Originally developed at Facebook, led by Simon Marlow
- Kind of like a scheduler for monadic computation
 - input "sequential", IO bound, monadic computation
 - will optimally¹ rewrite to parallelize
 - sophisticated caching
- Really great for time indexed databases where you really don't want to do more work than you have to.

[1] Uses `{-# LANGUAGE ApplicativeDo #-}`. Not actually optimal (couldn't be anyway), uses heuristics to avoid slow compile times and still get good results.

Example Computation

Suppose we have these functions:

```
-- | Core logic generated by QuasiQuoter,  
--   queries web3 api.  
balanceOf  
  :: MonadWeb3 m  
  => BlockNumber  
  -> Address  
  -> m (UIntN 256)  
...  
  
-- | Uses our stored history of all token transfers.  
getTraders  
  :: MonadPg m  
  => BlockNumber  
  -> Address  
  -> m [Address]  
...
```

Example Computation (*continued*)

Let **Neighbor** be a relations with **Neighbor(A,B)** if address A has traded with address B, where B currently has nonzero token balance.

```
-- | Gets the 10 Neighbors with the highest token
--   balances
getRichestNeighbors
  :: ( MonadWeb3 m
      , MonadPg m
      )
  => BlockNumber
  -> Address
  -> m [(Address, UIntN 256)]
getRichestNeighbors bn userAddress = do
  traders <- getTraders userAddress
  pairs <- forM traders $ \trader -> do
    bal <- balanceOf bn trader
    pure (trader, bal)
  let pairs' = filter ((> 0) . snd) pairs
  pure . take 10 . sortOn (Down . snd) $ pairs'
```

Example Computation (*continued*)

Introduce another parameter k with $\text{Neighbor}'(A, B, k)$ if there exists $A_1 \dots, A_{i-1}$ with $i \leq k$ and $\text{Neighbor}(A, A_1) \dots \text{Neighbor}(A_{i-1}, B)$.

In other words, A and B are at most k Neighbor pairs apart.

Example Computation: Naive Approach

getRichestNeighborsK

:: (MonadWeb3 m, MonadPg m)

=> BlockNumber

-> Int

-> Address

-> m [(Address, UIntN 256)]

getRichestNeighborsK bn k userAddress = **do**

traders <- getAllNs bn k userAddress

pure . take 10 . sortOn (Down . down) \$ pairs'

where

getAllNs _ 0 _ = pure []

getAllNs bn k userAddress = **do**

traders <- getTraders userAddress

pairs <- forM traders \$ \t -> **do**

bal <- balanceOf bn t

pure (t, bal)

let pairs' = filter ((> 0) . snd) pairs

addrs = map fst pairs'

rest <- concat <\$> mapM (getAllNs bn (k-1)) addrs

pure \$ pairs' <> rest

Example Computation: Problems

- `balanceOf` is possibly called multiple times per address.
- `getTraders` is possibly called multiple times per address.
- any time we used `mapM` (equivalently `forM`), we could be concurrently running each operation.

Bad Solutions

- Run computation in `State` monad holding `MVar` with maps
 - `traders :: Map Address [Address]`
 - `balances :: Map Address (UIntN 256)`
- Use combination of `par` strategies and `async / wait`.

Building custom concurrency like this is tedious, fragile, and not always composable.

Good Solutions

- Don't do anything at all.
- Let the compiler and libraries work for you.

Haxl Solution

The analogy is:

- Memory Management is to Garbage Collection as
Concurrency is to Haxl¹

You might be able to do better with custom concurrency controls, but unlikely. Also, it's probably not the actual problem you're trying to solve at the moment.

[1] Simon Marlow [StrangeLoop 2017](#)

Haxl Scheduler Diagram

Not an AST!

This represents a dependency graph where sequential dependencies are marked with $(>>=)$ and synchronization points with $(<*>)$.

